

EXTENDS *FiniteSets*, *Naturals*, *Sequences*

Constants: commands: set of records like $[key \mapsto \text{“key”}, value \mapsto \text{“value”}]$.
each command should be unique in the set.

replicas: set of replicas.

CONSTANTS *commands*, *replicas*

Variables: leader: the current leader. epoch: the current epoch (the number of leader changes).

specPools: the spec pool of each replica.

requested: the set of requested commands (by client).

committed: the set of committed commands by *CURP*. It also records the index of the last same-key command in the *syncd* sequence at the time.

unsyncd: the sequence of *unsyncd* commands. *syncd*: the sequence of *syncd* commands.

VARIABLES *leader*, *epoch*, *specPools*, *requested*, *committed*, *unsyncd*, *syncd*

The initial state of the system.

Init \triangleq
 $\wedge leader \in replicas$
 $\wedge epoch = 1$
 $\wedge specPools = [r \in replicas \mapsto \{\}]$
 $\wedge requested = \{\}$
 $\wedge committed = \{\}$
 $\wedge unsyncd = \langle \rangle$
 $\wedge syncd = \langle \rangle$

Helper function for converting a set to a set containing all sequences containing the elements of the set exactly once and no other elements.

SetToSeqs(*set*) \triangleq
 LET *len* $\triangleq 1 \dots Cardinality(set)$ IN
 $\{f \in [len \rightarrow set] : \forall i, j \in len : i \neq j \Rightarrow f[i] \neq f[j]\}$

Helper function for getting the index of the last element in a sequence satisfying the predicate.

GetIdxInSeq(*seq*, *Test*($_$)) \triangleq
 LET *I* $\triangleq \{i \in 1 \dots Len(seq) : Test(seq[i])\}$ IN
 IF *I* $\neq \{\}$ THEN CHOOSE *i* $\in I : \forall j \in I : j \leq i$ ELSE 0

SuperQuorum: In $N = 2 * f + 1$ replicas, a *SuperQuorum* is a set of replicas that contains at least $f + (f + 1) / 2 + 1$ replicas.

The client can consider a command as committed if and only if it receives positive responses from a set of replicas larger then a *SuperQuorum*.

IsSuperQuorum(*S*) \triangleq
 LET *f* $\triangleq (Cardinality(replicas) - 1) \div 2$
 size $\triangleq f + (f + 1) \div 2 + 1$
 IN $Cardinality(S) \geq size$

LeastQuorum: In $N = 2 * f + 1$ replicas, a *LeastQuorum* is a set of replicas that contains at least $(f + 1) / 2 + 1$ replicas.

When a replica becomes a leader, it must recover the command if and only if the command is in the *specPool* of a set of replicas larger then a *LeastQuorum*.

$IsLeastQuorum(S) \triangleq$
 LET $f \triangleq (Cardinality(replicas) - 1) \div 2$
 $size \triangleq (f + 1) \div 2 + 1$
 IN $Cardinality(S) \geq size$

RecoveryQuorums: In $N = 2 * f + 1$ replicas, a *RecoveryQuorum* is a set of replicas that contains at least $f + 1$ replicas.

A replica must gather a *RecoveryQuorum* of *replicas'* *specPool* to recover the commands that need to be recovered.

This defines the set consisting of all *RecoveryQuorums*.

$recoveryQuorums \triangleq$
 LET $f \triangleq (Cardinality(replicas) - 1) \div 2$
 $size \triangleq f + 1$
 IN $\{q \in \text{SUBSET } replicas : Cardinality(q) = size\}$

The Abstraction of the normal procedure of *CURP*.

$Request \triangleq$
 $\exists cmd \in commands \setminus requested :$
 $\wedge requested' = requested \cup \{cmd\}$

To simulate the unreliability of the network,
 only a subset of replicas could receive the request.

$\wedge \exists received \in \text{SUBSET } replicas :$

The set of replicas that got no conflict in the spec pool.

$\wedge \text{LET } acceptedReplicas \triangleq$
 $\{r \in received :$
 $\quad \forall specCmd \in specPools[r] :$
 $\quad \quad specCmd.key \neq cmd.key\}$

IN

Update the *specPool*.

$\wedge specPools' = [r \in replicas \mapsto$
 IF $r \in acceptedReplicas$
 THEN $specPools[r] \cup \{cmd\}$
 ELSE $specPools[r]$

If there is at least a superquorum set of replicas that accepted
 the request, and the leader can execute the command,
 the request is committed.

$\wedge \text{LET } CompareKey(elem) \triangleq elem.key = cmd.key$ IN
 IF
 $\wedge IsSuperQuorum(acceptedReplicas)$

$\wedge leader \in acceptedReplicas$
 $\wedge GetIdxInSeq(unsynced, CompareKey) = 0$
 THEN
 The previous state of the key is also recorded.
 This is used to check the correctness of the property.
 LET $prevIdx \triangleq GetIdxInSeq(synced, CompareKey)$ IN
 $committed' = committed \cup \{[$
 $cmd \mapsto cmd,$
 $prevIdx \mapsto prevIdx]\}$
 ELSE $committed' = committed$

No matter if the request is committed or not,
 as long as the leader is in the received set,
 the command should be *synced* afterward.
 \wedge IF $leader \in received$
 THEN $unsynced' = Append(unsynced, cmd)$
 ELSE $unsynced' = unsynced$

\wedge UNCHANGED $\langle leader, epoch, synced \rangle$

Syncing a command using the back-end protocol like Raft. The implementation details of the back-end protocol are omitted.

$Sync \triangleq$
 $\wedge unsynced \neq \langle \rangle$
 $\wedge specPools' = [r \in replicas \mapsto specPools[r] \setminus \{Head(unsynced)\}]$
 $\wedge synced' = Append(synced, Head(unsynced))$
 $\wedge unsynced' = Tail(unsynced)$
 \wedge UNCHANGED $\langle leader, epoch, requested, committed \rangle$

Leader Change Action

The new leader should gather at least a *RecoveryQuorum* of *replicas'* *specPool* to recover the commands.

Commands existed in the *specPool* of a *LeastQuorum* of replicas need to be recovered.

$LeaderChange \triangleq$
 $\exists newLeader \in (replicas \setminus \{leader\}) :$
 $\wedge leader' = newLeader$
 $\wedge epoch' = epoch + 1$
 $\wedge \exists recoveryQuorum \in recoveryQuorums :$
 LET $specCmds \triangleq \text{UNION } \{specPools[r] : r \in recoveryQuorum\}$
 $newSpecPool \triangleq \{cmd \in specCmds : IsLeastQuorum(\{r \in replicas : cmd \in specPools[r]\})\}$
 IN
 $\wedge specPools' = [specPools \text{ EXCEPT } ![newLeader] = newSpecPool]$
 $\wedge unsynced' \in SetToSeqs(newSpecPool)$
 \wedge UNCHANGED $\langle requested, committed, synced \rangle$

$$\begin{aligned}
Next &\triangleq \\
&\vee Request \\
&\vee Sync \\
&\vee LeaderChange
\end{aligned}$$

$$Spec \triangleq Init \wedge \Box[Next]_{\langle leader, epoch, specPools, requested, committed, unsynced, synced \rangle}$$

Type Check

$$\begin{aligned}
TypeOK &\triangleq \\
&\wedge leader \in replicas \\
&\wedge epoch \in Nat \\
&\wedge \forall r \in replicas : specPools[r] \subseteq commands \\
&\wedge requested \subseteq commands \\
&\wedge \forall committedCmd \in committed : \\
&\quad \wedge committedCmd.cmd \in commands \\
&\quad \wedge committedCmd.prevIdx \in 0 \dots Len(synced) \\
&\wedge synced \in \{SetToSeqs(s) : s \in SUBSET\ commands\} \\
&\wedge unsynced \in \{SetToSeqs(s) : s \in SUBSET\ commands\}
\end{aligned}$$

Stability Property

This is the key property of *CURP*. There are two parts of the property.

1. If a command is committed by *CURP*, command will eventually be *synced* by the back-end protocol.
2. If a command is committed by *CURP*, when the command is *synced* by the back-end protocol, there will never be a command with the same key between the command and the recorded previous same-key command in the *synced* sequence.

$$\begin{aligned}
Stability &\triangleq \\
&\forall committedCmd \in committed : \\
&\quad LET\ CompareExact(elem) \triangleq elem = committedCmd.cmd \\
&\quad\quad syncedIdx \triangleq GetIdxInSeq(synced, CompareExact) \\
&\quad IN \\
&\quad \wedge syncedIdx \neq 0 \\
&\quad \wedge \forall j \in (committedCmd.prevIdx + 1) \dots (syncedIdx - 1) : \\
&\quad\quad synced[j].key \neq committedCmd.cmd.key
\end{aligned}$$

THEOREM $Spec \Rightarrow \Box TypeOK \wedge \Diamond Stability$