MODULE *curp*

EXTENDS *FiniteSets*, *Naturals*, *Sequences*

Constants:
commands : *set of records like*[*key* $\mapsto$ "key", *value* $\mapsto$ "value"].
replicas : *set of replicas*.

CONSTANTS *commands*, *replicas*

ASSUME *IsFiniteSet*(*replicas*)

Variables :
leader : *records the leader of each epoch*.
epoch : *current epoch*(*the number of leader changes*).
proposedCmds : *the set of proposed commands*.
proposeRequests : *the set of propose could be received by each replica*.
proposeResponses : *the set of responses in each epoch to each proposed command*.
specPools : *the speculative pool of each replica*.
uncommittedCmds : *the sequence of back* $-$ *end protocol uncommitted commands*.
committedCmds : *the sequence of back* $-$ *end protocol committed commands*.
commitMsgs : *the set of commit messages could be received by each replica*.
specExecPrevCmd : *the index of the last same* $-$ *key command in the committed sequence at the time the leader responds to the proposal*.

VARIABLES *leader*, *epoch*, *proposedCmds*, *proposeRequests*, *proposeResponses*, *specPools*, *uncommittedCmds*, *committedCmds*, *commitMsgs*, *specExecPrevCmd*

The epoch space.

$epoches \triangleq Nat$

Special noLeader value, for future epoches.

$noLeader \triangleq \text{CHOOSE } r : r \notin replicas$

In $N = 2 * f + 1$ *replicas* :
quorum : *a set of replicas that contains at least* $f + 1$ *replicas*.
superQuorum : *a set of replicas that contains at least* $f + (f + 1)/2 + 1$ *replicas*.
recoverQuorum : *a set of replicas that contains at least*$(f + 1)/2 + 1$ *replicas*.

$$
\begin{aligned}
&quorums \triangleq \\
&\quad \text{LET } f \triangleq Cardinality(replicas) \div 2 \\
&\qquad size \triangleq f + 1 \\
&\quad \text{IN } \{q \in \text{SUBSET } replicas : Cardinality(q) \geq size\}
\end{aligned}
$$

$$
\begin{aligned}
&superQuorums \triangleq \\
&\quad \text{LET } f \triangleq Cardinality(replicas) \div 2 \\
&\qquad size \triangleq f + (f + 1) \div 2 + 1
\end{aligned}
$$

```
    IN   {q ∈ SUBSET replicas : Cardinality(q) ≥ size}

recoverQuorums ≜
    LET f ≜ Cardinality(replicas) ÷ 2
        size ≜ (f + 1) ÷ 2 + 1
    IN   {q ∈ SUBSET replicas : Cardinality(q) ≥ size}

Helper function for converting a set to a set containing all sequences
containing the elements of the set exactly once and no other elements.
SetToSeqs(set) ≜
    LET len ≜ 1 .. Cardinality(set)IN
        {f ∈ [len → set] : ∀ i, j ∈ len : i ≠ j ⇒ f[i] ≠ f[j]}

Helper function for getting the index of the last element in a sequence
satisfying the predicate.
GetIdxInSeq(seq, Pred(_)) ≜
    LET I ≜ {i ∈ 1 .. Len(seq) : Pred(seq[i])}IN
        IF I ≠ {} THEN CHOOSE i ∈ I : ∀ j ∈ I : j ≤ i ELSE 0

Propose a command.
This is done by the client sending a proposeRequest to all replicas.
Propose(cmd) ≜
    ∧ proposedCmds′ = proposedCmds ∪ {cmd}
    ∧ proposeRequests′ =
        [r ∈ replicas ↦ proposeRequests[r] ∪ {cmd}]
    ∧ UNCHANGED ⟨leader, epoch, specPools, proposeResponses,
                 uncommittedCmds, committedCmds, commitMsgs,
                 specExecPrevCmd⟩

How the leader process a proposeRequest.
ProcessProposeLeader(r, cmd) ≜
    LET specPoolHasConflict ≜
            ∃ specCmd ∈ specPools[r] : specCmd.key = cmd.key
        uncommittedCmdsHasConflict ≜
            GetIdxInSeq(uncommittedCmds, LAMBDA e : e.key = cmd.key) ≠ 0
    IN
        ∧ proposeRequests′ =
            [proposeRequests EXCEPT ![r] = @ \ {cmd}]
        ∧ specPools′ =
            [specPools EXCEPT ![r] =
                IF ¬specPoolHasConflict THEN @ ∪ {cmd} ELSE @]
        ∧ uncommittedCmds′ = Append(uncommittedCmds, cmd)
        ∧ proposeResponses′ =
            [proposeResponses EXCEPT ![cmd][epoch] =
                IF ¬specPoolHasConflict ∧ ¬uncommittedCmdsHasConflict
```

```
                THEN @ ∪ {r}
                ELSE @]
        ∧ specExecPrevCmd′ =
            [specExecPrevCmd EXCEPT ![cmd] =
                IF ¬specPoolHasConflict ∧ ¬uncommittedCmdsHasConflict
                THEN GetIdxInSeq(committedCmds, LAMBDA e : e.key = cmd.key)
                ELSE @]
        ∧ UNCHANGED ⟨leader, epoch, proposedCmds, committedCmds,
                     commitMsgs⟩
```

How a non − leader replica process a proposeRequest.

```
ProcessProposeNonLeader(r, cmd) ≜
    LET specPoolHasConflict ≜
            ∃ specCmd ∈ specPools[r] : specCmd.key = cmd.key
    IN
        ∧ proposeRequests′ =
            [proposeRequests EXCEPT ![r] = @ \ {cmd}]
        ∧ specPools′ =
            [specPools EXCEPT ![r] =
                IF ¬specPoolHasConflict THEN @ ∪ {cmd} ELSE @]
        ∧ proposeResponses′ =
            [proposeResponses EXCEPT ![cmd][epoch] =
                IF ¬specPoolHasConflict THEN @ ∪ {r} ELSE @]
        ∧ UNCHANGED ⟨leader, epoch, proposedCmds, uncommittedCmds,
                     committedCmds, commitMsgs, specExecPrevCmd⟩
```

Syncing a command using the back − end protocol(Raft).The implementation details are omitted.

A replica may not be able to receive the commit message at the exact time the the leader sends it.

```
Commit ≜
    ∧  committedCmds′ = Append(committedCmds, Head(uncommittedCmds))
    ∧  commitMsgs′ =
         [r ∈ replicas ↦ commitMsgs[r] ∪ {Head(uncommittedCmds)}]
    ∧  uncommittedCmds′ = Tail(uncommittedCmds)
    ∧  UNCHANGED ⟨leader, epoch, specPools, proposedCmds, proposeRequests,
                  proposeResponses, specExecPrevCmd⟩
```

How a replica process a commit message.

```
ProcessCommitMsg(r, cmd) ≜
    ∧ commitMsgs′ =
        [commitMsgs EXCEPT ![r] = @ \ {cmd}]
    ∧ specPools′ = [specPools EXCEPT ![r] = @ \ {cmd}]
    ∧ UNCHANGED ⟨leader, epoch, proposedCmds, proposeRequests,
                 proposeResponses, uncommittedCmds, committedCmds,
```

$specExecPrevCmd\rangle$

Leader Change Action

The new leader should gather at least a quorum of replicas specPool to recover the commands.

Commands existed in the *specPool* of a *RecoverQuorum* of replicas need to be recovered.

$$
\begin{array}{l}
LeaderChange(l) \triangleq \\
\quad \land leader' = [e \in epoches \mapsto \text{IF } e = epoch + 1 \text{ THEN } l \text{ ELSE } leader[e]] \\
\quad \land epoch' = epoch + 1 \\
\quad \land \exists\, q \in quorums : \\
\qquad \text{LET } specCmds \triangleq \text{UNION } \{specPools[r] : r \in q\} \\
\qquad\quad newSpecPool \triangleq \\
\qquad\qquad \{cmd \in specCmds : \\
\qquad\qquad\quad \{r \in q : cmd \in specPools[r]\} \in recoverQuorums\} \\
\qquad \text{IN} \\
\qquad\quad \land specPools' = [specPools \text{ EXCEPT } ![l] = newSpecPool] \\
\qquad\quad \land uncommittedCmds' \in SetToSeqs(newSpecPool) \\
\quad \land \text{UNCHANGED } \langle proposedCmds,\ proposeRequests,\ proposeResponses, \\
\qquad\qquad committedCmds,\ commitMsgs,\ specExecPrevCmd\rangle
\end{array}
$$

The initial state of the system.

$$
\begin{array}{l}
Init \triangleq \\
\quad \exists\, r \in replicas : \\
\qquad \text{LET } initEpoch \triangleq 1 initLeader \triangleq r \text{IN} \\
\qquad\quad \land leader = [e \in epoches \mapsto \\
\qquad\qquad \text{IF } e = initEpoch \text{ THEN } initLeader \text{ ELSE } noLeader] \\
\qquad\quad \land epoch = initEpoch \\
\qquad\quad \land proposedCmds = \{\} \\
\qquad\quad \land proposeRequests = [replica \in replicas \mapsto \{\}] \\
\qquad\quad \land proposeResponses = \\
\qquad\qquad [cmd \in commands \mapsto [e \in epoches \mapsto \{\}]] \\
\qquad\quad \land specPools = [replica \in replicas \mapsto \{\}] \\
\qquad\quad \land uncommittedCmds = \langle\rangle \\
\qquad\quad \land committedCmds = \langle\rangle \\
\qquad\quad \land commitMsgs = [replica \in replicas \mapsto \{\}] \\
\qquad\quad \land specExecPrevCmd = [cmd \in commands \mapsto 0]
\end{array}
$$

$$
\begin{array}{l}
Next \triangleq \\
\quad \lor \exists\, cmd \in (commands \setminus proposedCmds) : Propose(cmd) \\
\quad \lor \exists\, r \in replicas : \exists\, cmd \in proposeRequests[r] : \\
\qquad \text{IF } leader[epoch] = r \\
\qquad\quad \text{THEN } ProcessProposeLeader(r,\ cmd) \\
\qquad\quad \text{ELSE } ProcessProposeNonLeader(r,\ cmd) \\
\quad \lor uncommittedCmds \neq \langle\rangle \land Commit \\
\quad \lor \exists\, r \in replicas : \exists\, cmd \in commitMsgs[r] : ProcessCommitMsg(r,\ cmd) \\
\quad \lor \exists\, l \in replicas : LeaderChange(l)
\end{array}
$$

$$Spec \triangleq Init \land \Box[Next]\langle leader,\ epoch,\ specPools,\ proposedCmds,\ proposeRequests,\ proposeResponses,\ uncommittedCmds,\ committedCmds,\ commitMsgs,\ specExecPrevCmd\rangle$$

Type Invariants

$$\begin{aligned}
&TypeOK \triangleq \\
&\quad \land \quad leader \in [epoches \to (replicas \cup \{noLeader\})] \\
&\quad \land \quad epoch \in epoches \\
&\quad \land \quad proposedCmds \subseteq commands \\
&\quad \land \quad proposeRequests \in [replicas \to \text{SUBSET } commands] \\
&\quad \land \quad proposeResponses \in [commands \to [epoches \to \text{SUBSET } replicas]] \\
&\quad \land \quad specPools \in [replicas \to \text{SUBSET } commands] \\
&\quad \land \quad uncommittedCmds \in \text{UNION } \{SetToSeqs(s) : s \in \text{SUBSET } commands\} \\
&\quad \land \quad committedCmds \in \text{UNION } \{SetToSeqs(s) : s \in \text{SUBSET } commands\} \\
&\quad \land \quad commitMsgs \in [replicas \to \text{SUBSET } commands] \\
&\quad \land \quad specExecPrevCmd \in [commands \to 0 \,..\, Cardinality(commands)]
\end{aligned}$$

Stability Property

This is the key property of *CURP*:

1. If a command is committed by *CURP*, command will eventually be synced by the back-end protocol.
2. If a command is committed by *CURP*, when the command is synced be the back-end protocol, there will never be a command with the same key between the command and the recorded previous same-key command in the synced sequence.

$$\begin{aligned}
&Stability \triangleq \\
&\quad \forall\, cmd \in commands : \forall\, e \in epoches : \\
&\qquad (\land\, leader[e] \in proposeResponses[cmd][e] \\
&\qquad\ \land\, proposeResponses[cmd][e] \in superQuorums) \Rightarrow \\
&\qquad\quad \text{LET } idx \triangleq GetIdxInSeq(committedCmds, \text{ LAMBDA } t : t = cmd) \\
&\qquad\qquad\ \ prevExecCmds \triangleq SubSeq(committedCmds, 1, idx) \\
&\qquad\quad \text{IN} \\
&\qquad\qquad \land\, idx \neq 0 \\
&\qquad\qquad \land\, GetIdxInSeq(prevExecCmds, \text{ LAMBDA } t : t.key = cmd.key) = \\
&\qquad\qquad\qquad specExecPrevCmd[cmd]
\end{aligned}$$

$$\text{THEOREM } Spec \Rightarrow \Box\, TypeOK \land \Diamond Stability$$